



Containerization: Best Practices & Advanced Topics

Peter Z. Vaillancourt

Computational Scientist

Center for Advanced Computing (CAC)

Cornell University

XCRI Engineer, XSEDE

XSEDE

Extreme Science and Engineering
Discovery Environment



Containerization: Best Practices & Advanced Topics

Outline

- Lifecycle of a Container
 - Development vs. Production
 - Reducing Container Sizes
 - Data Management
- Deploying a Container
 - Registries
 - In the Cloud
 - On a Shared Resource
- Security
- Next Steps
 - Reproducible Containers
 - Container Orchestration





Lifecycle of a Container

Containers for Development vs. Production

1. Production: Containers as a software distribution method


- Portability of a consistent environment for users
- Easily distributed
- Highly accessible
- Pre-packaged software containers often require customization



Lifecycle of a Container

Containers for Development vs. Production

1. **Production:** Containers as a software distribution method
 - Portability of a consistent environment for users
 - Easily distributed
 - Highly accessible
 - Pre-packaged software containers often require customization
2. **Development:** Containers as a development environment
 - Builds a consistent environment early, including dependencies
 - Useful for teams of developers/researchers
 - Larger if including dev tools
 - Often requires cleanup for production



Lifecycle of a Container

Containers for Development vs. Production

Development

Production



Lifecycle of a Container

Containers for Development vs. Production

Development

- Contains dependencies, code, environment variables, etc.
- No real size limit: text editors, VNC, data visualization, etc.
- Code is changed and updated
- Runs can be varied and versatile to initiate

Production



Lifecycle of a Container

Containers for Development vs. Production

Development

- Contains dependencies, code, environment variables, etc.
- No real size limit: text editors, VNC, data visualization, etc.
- Code is changed and updated
- Runs can be varied and versatile to initiate

Production


- Contains dependencies, code, environment variables, etc.
- Should be as lightweight as possible: no need for nice aesthetic features
- Code is static
- Requires a run script or easy commands




Lifecycle of a Container

Reducing Container Sizes

- Docker Layers
 - Base image
 - CentOS 215MB
 - Debian 114MB
 - Ubuntu 73.9MB
 - Alpine 5.57MB







Lifecycle of a Container

Reducing Container Sizes

- Docker Layers
 - Base image
 - CentOS 215MB
 - Debian 114MB
 - Ubuntu 73.9MB
 - Alpine 5.57MB
 - Certain commands add layers:
RUN, ADD, COPY
 - 1 instruction = 1 layer
 - Other commands create temporary layers
 - Also see the [Docker docs](#)





Lifecycle of a Container

Reducing Container Sizes

- Combining multiple commands
 - Pip commands can use a requirements file

Example requirements.txt:

```
alembic
fitsio==0.9.11
requests_oauthlib
marshmallow
ephem
scikit-sparse
corner
numexpr
astropy
runipy
...
```



Lifecycle of a Container

Reducing Container Sizes

- Combining multiple commands
 - Pip commands can use a requirements file
 - If using several RUN commands in a row, it's an opportunity to combine:

```
RUN wget -q https://bitbucket.org/psrsoft/tempo2/get/master.tar.gz && \  
tar xzf master.tar.gz && \  
cd psrsoft-tempo2-* && \  
./bootstrap && \  
CPPFLAGS="-I/opt/pulsar/include" LDFLAGS="-L/opt/pulsar/lib" ./configure -- prefix=/opt/pulsar --with-calceph=/opt/pulsar && \  
make && make install && make plugins && make plugins-install && \  
mkdir -p /opt/pulsar/share/tempo2 && \  
cp -Rp T2runtime/* /opt/pulsar/share/tempo2/. && \  
cd .. && rm -rf psrsoft-tempo2-* master.tar.gz
```



Lifecycle of a Container

Reducing Container Sizes

- Combining multiple commands
 - Pip commands can use a requirements file
 - If using several RUN commands in a row, it's an opportunity to combine
- Use multi-stage builds
 - Leverages docker build cache
 - Can be used to remove some contents after build for security



Lifecycle of a Container

Reducing Container Sizes

- Combining multiple commands
 - Pip commands can use a requirements file
 - If using several RUN commands in a row, it's an opportunity to combine
- Use multi-stage builds
 - Leverages docker build cache
- Don't install what you don't need
- Multiple decoupled containers (microservices)



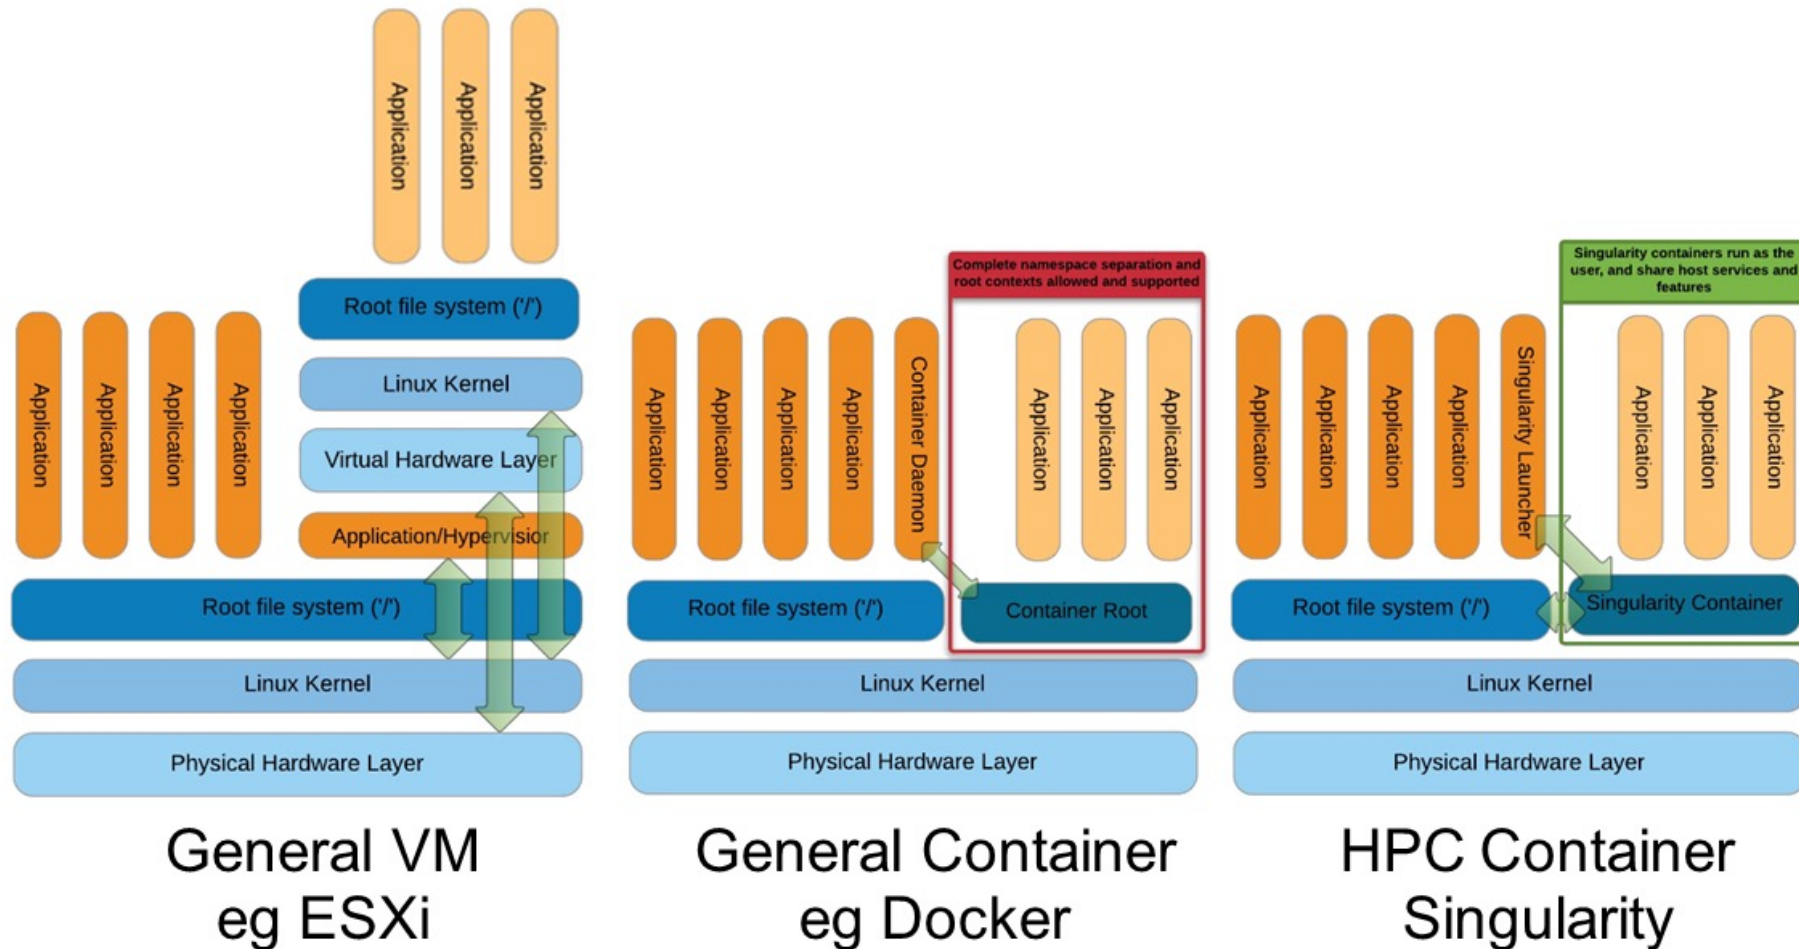
Lifecycle of a Container

Reducing Container Sizes

- Using a prebuilt image from a registry?
 - Decide whether to build your own
 - Look at the layers to see what is excess
 - Build files on GitHub can reveal unnecessary additions
- Building your own is the best way to ensure you have only what you need



Lifecycle of a Container Data Management






Lifecycle of a Container

Data Management


- Manage data in Docker:
 - <https://docs.docker.com/storage/>
 - Volumes
 - <https://docs.docker.com/storage/volumes/>
 - Bind mounts
 - <https://docs.docker.com/storage/bind-mounts/>
 - tmpfs mounts
 - <https://docs.docker.com/storage/tmpfs/>
- Manage data in Singularity
 - Access to the root filesystem:
 - https://sylabs.io/guides/3.5/user-guide/quick_start.html#working-with-files
 - Bind Paths and Mounts:
 - https://sylabs.io/guides/3.5/user-guide/bind_paths_and_mounts.html
 - Persistent Overlays:
 - https://sylabs.io/guides/3.5/user-guide/persistent_overlays.html



Deploying a Container

Uploading Containers

- Containers on public registries can be a great way to
 - Share scientific software portably
 - Help others reproduce your research
 - Create a community around a software/workflow
 - Set up automated builds (connect your GitHub repo)
- Common Public Registries
 - DockerHub
 - Sylabs Cloud Hosting (11GB limit)




Deploying a Container

Best Practices for Uploading Containers


- **Don't upload**
 - Private data – very important for research
 - Private or licensed software
- **Do include**
 - Software licenses
 - Documentation
 - Software and dependencies
 - Runscripts for production



Deploying a Container In the Cloud




- Will it work in the cloud?
 - Moving from HPC adds complexity
 - MPI
 - May require container orchestration
 - Data management
- Use Docker
 - Public cloud providers offer managed services
 - Container Orchestration options
 - Ease of use







Deploying a Container In the Cloud



- What is the right cloud?
 - Private clouds using OpenStack have various options
 - Public clouds have vendor-specific options galore
 - Multi-cloud or generalized tools also available
- Security is a consideration






Deploying a Container On HPC Resources



- Containers can simplify getting started
 - No need to install to your home directory
 - No need to pester sysadmins to install your software
- Using Singularity on XSEDE
 - It's available and secure
 - Automatic mounts for easy data access
 - Multinode orchestration via Slurm
- Job scripts and bind mounts may vary on different systems





Deploying a Container On HPC Resources



- Singularity with MPI
 - MPI **major version** in the container *must* match the host
<https://sylabs.io/guides/3.5/user-guide/mpi.html>
 - Singlenode will work if MPI is run inside the container
 - Multinode is more dependent on system
 - Choice of MPI library may be limited by hardware
 - Scheduler integrations, such as Slurm





Security

Root Access

- Sensitive systems cannot allow root access
- XSEDE systems use Singularity
- Another option is Docker Rootless Mode
 - [Docker Docs on Rootless Mode](#)
 - [DockerCon 2020 Talk on Rootless Mode](#)
 - Recently became a [fully supported feature](#)
- Setup a user or users for shared Docker containers (same as shared system)
https://docs.docker.com/develop/develop-images/dockerfile_best-practices



Security Cloud VMs


- Implement security at a Virtual Machine (VM) level
 - Firewall
 - Security Groups
 - Limit ssh access
- For public images
 - Pay attention to what they contain
 - Look for the Dockerfile
 - GitHub repo



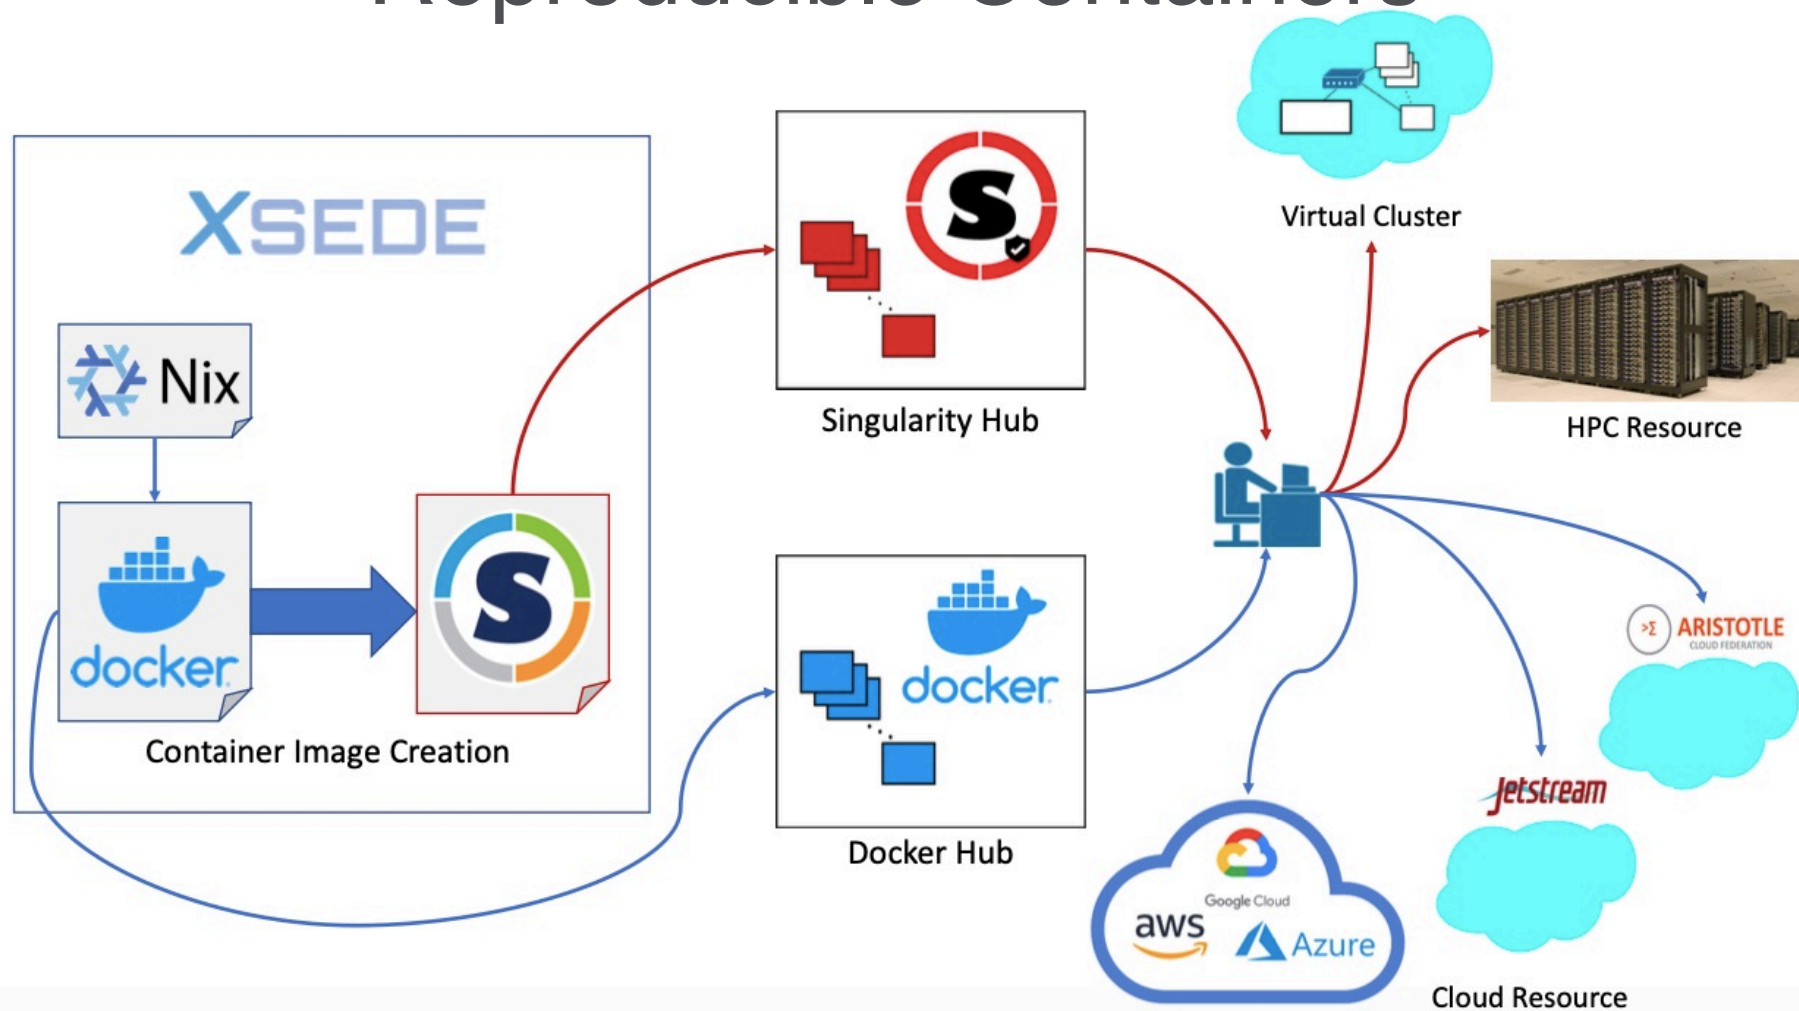
Security


Other Security Features

- Docker vulnerability scanning
 - <https://docs.docker.com/engine/scan>
 - <https://github.com/docker/scan-cli-plugin>
- Singularity Security Options
 - Linux Capabilities
 - Encrypted Containers
 - https://sylabs.io/guides/3.5/user-guide/security_options.html#security-options





Next Steps Reproducible Containers





Next Steps

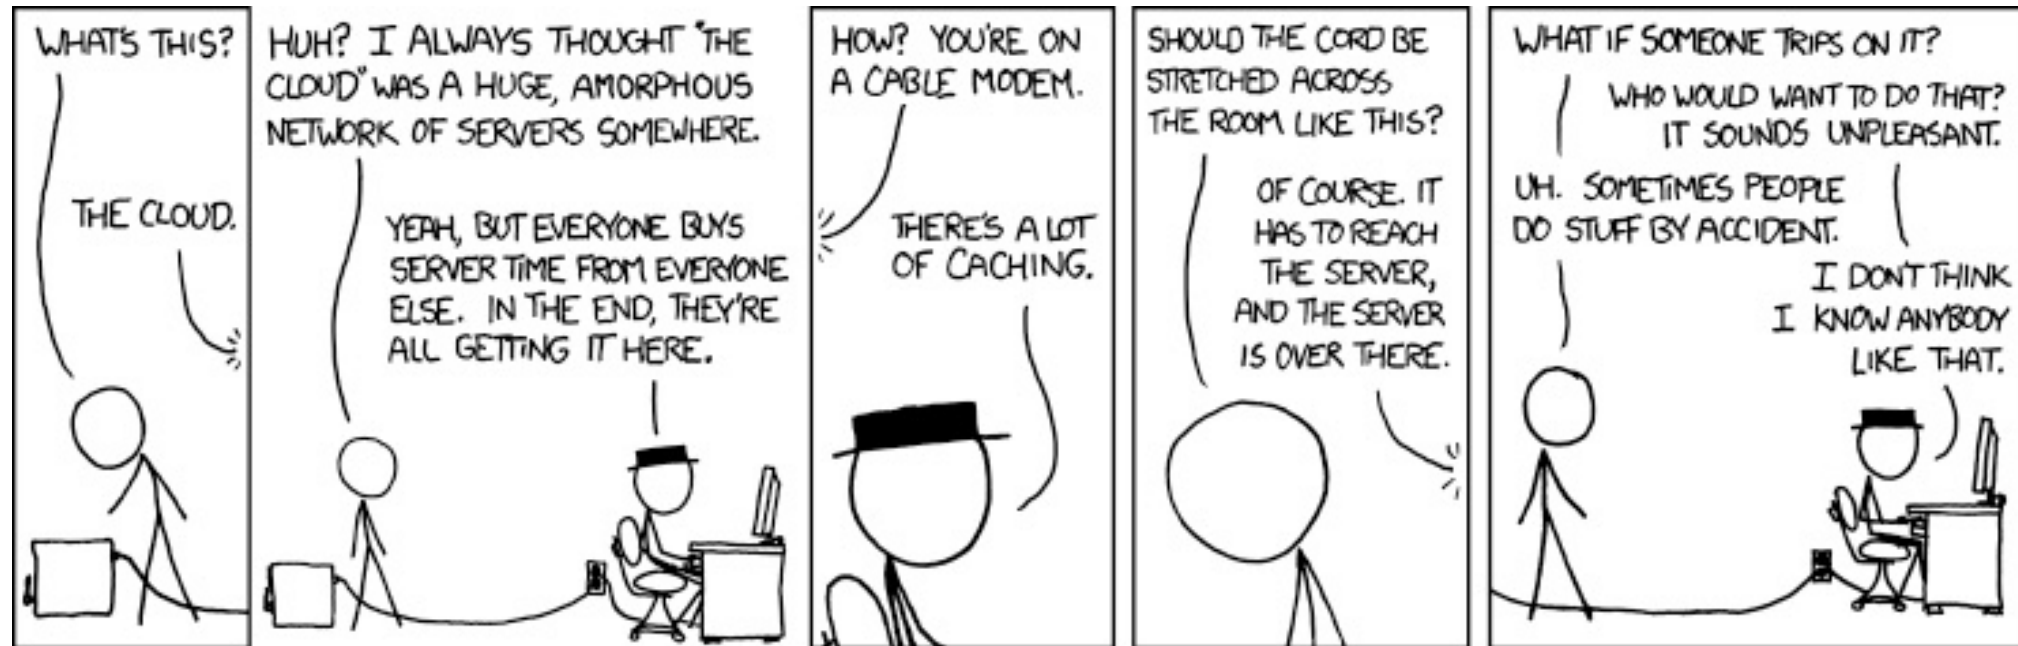
Container Orchestration




+ ANSIBLE



Questions?





XSEDE

Extreme Science and Engineering
Discovery Environment

Thank you!


https://github.com/XSEDE/Container_Tutorial

Peter Vaillancourt
Computational Scientist
Center for Advanced Computing (CAC)
Cornell University
XCRI Engineer, XSEDE



Lifecycle of a Container

Example Container: Radio Astronomy






Lifecycle of a Container

Example Container: Radio Astronomy

- Started with a NANOGrav container: [nanograv/nanopulsar](#)
 - Based on [jupyter/datascience-notebooks](#) (includes Python, R, and more)
 - Wide variety of Radio Astronomy software and tools






Lifecycle of a Container

Example Container: Radio Astronomy

- Started with a NANOGrav container: [nanograv/nanopulsar](#)
 - Based on [jupyter/datascience-notebooks](#) (includes Python, R, and more)
 - Wide variety of Radio Astronomy software and tools
- After 1 year, needed to be updated: [federatedcloud/nanopulsar](#)
- Used for development with additions: [federatedcloud/modulation_index](#)
 - ~11GB for just dependencies





Lifecycle of a Container

Example Container: Radio Astronomy

- Started with a NANOGrav container: [nanograv/nanopulsar](#)
 - Based on [jupyter/datascience-notebooks](#) (includes Python, R, and more)
 - Wide variety of Radio Astronomy software and tools
- After 1 year, needed to be updated: [federatedcloud/nanopulsar](#)
- Used for development with additions: [federatedcloud/modulation_index](#)
 - ~11GB for just dependencies
- Created a minimal container for production runs
 - ~3GB for just dependencies
 - Docker version: [federatedcloud/docker-PRESTO](#)
 - Singularity version: [federatedcloud/singularity-PRESTO](#)

