

Cornell University

Containerization: Best Practices & Advanced Topics

Center for Advanced Computing (CAC)
Cornell University
XSEDE

XSEDE

Extreme Science and Engineering
Discovery Environment

Containerization: Best Practices & Advanced Topics

Outline

- Lifecycle of a Container
 - Development vs Production
 - Reducing Container Sizes
 - Data Management
- Deploying a Container
 - Registries
 - In the Cloud
 - On a Shared Resource
- Security
- Next Steps
 - Reproducible Containers
 - Container Orchestration



Lifecycle of a Container

Containers for Development vs. Production

Production: Containers as a software distribution method

- Portability of a consistent environment for users
- Easily distributed
- Highly accessible
- Pre-packaged software containers often require customization

Lifecycle of a Container

Containers for Development vs. Production

Production: Containers as a software distribution method

- Portability of a consistent environment for users
- Easily distributed
- Highly accessible
- Pre-packaged software containers often require customization

Development: Containers as a development environment

- Builds a consistent environment early, including dependencies
- Useful for teams of developers/researchers
- Larger if including dev tools
- Often requires cleanup for production

Lifecycle of a Container

Containers for Development vs. Production

Development

- Contains dependencies, code, environment variables, etc
- No real size limit: text editors, VNC, data visualization, etc
- Code is changed and updated
- Runs can be varied and versatile to initiate

Production


- Contains dependencies, code, environment variables, etc.
- Should be as lightweight as possible: no need for nice aesthetic features
- Code is static
- Requires a run script or easy commands

Lifecycle of a Container

Reducing Container Sizes

Docker Layers

- Base image
 - Rocky 205MB
 - Alma 189 MB
 - Debian 114MB
 - Ubuntu 73.9MB
 - Alpine 5.57MB




Lifecycle of a Container

Reducing Container Sizes

Docker Layers

- Base image
 - Rocky 205MB
 - Alma 189MB
 - Debian 114MB
 - Ubuntu 73.9MB
 - Alpine 5.57MB
- Certain commands add layers:
RUN, ADD, COPY
- 1 instruction = 1 layer
- Other commands create temporary layers
- Also see the Docker docs



Lifecycle of a Container

Reducing Container Sizes

Combining multiple commands

- Pip commands can use a requirements file

Example requirements.txt:

```
alembic
fitsio==0.9.11
requests_oauthlib
marshmallow
ephem
scikit-sparse
corner
numexpr
astropy
runipy
...
```


Lifecycle of a Container

Reducing Container Sizes

Combining multiple commands

- Pip commands can use a requirements file
- If using several RUN commands in a row, this is an opportunity to combine:

```
RUN wget -q https://bitbucket.org/psrsoft/tempo2/get/master.tar.gz && \  
tar xzf master.tar.gz && \  
cd psrsoft-tempo2-* && \  
./bootstrap && \  
CPPFLAGS="-I/opt/pulsar/include" LDFLAGS="-L/opt/pulsar/lib" \  
./configure --prefix=/opt/pulsar --with-calceph=/opt/pulsar && \  
make && make install && make plugins && make plugins-install && \  
mkdir -p /opt/pulsar/share/tempo2 && \  
cp -Rp T2runtime/* /opt/pulsar/share/tempo2/. && \  
cd .. && rm -rf psrsoft-tempo2-* master.tar.gz
```

Lifecycle of a Container

Reducing Container Sizes

Combining multiple commands

- Pip commands can use a requirements file
- If using several RUN commands in a row, this is an opportunity to combine:

Use multi-stage builds

- Leverages docker build cache
- Can be used to remove some contents after the build for security

Lifecycle of a Container

Reducing Container Sizes

Combining multiple commands

- Pip commands can use a requirements file
- If using several RUN commands in a row, this is an opportunity to combine:

Use multi-stage builds

- Leverages docker build cache
- Can be used to remove some contents after the build for security

Don't install what you don't need

Multiple decoupled containers (microservices)

Lifecycle of a Container

Reducing Container Sizes

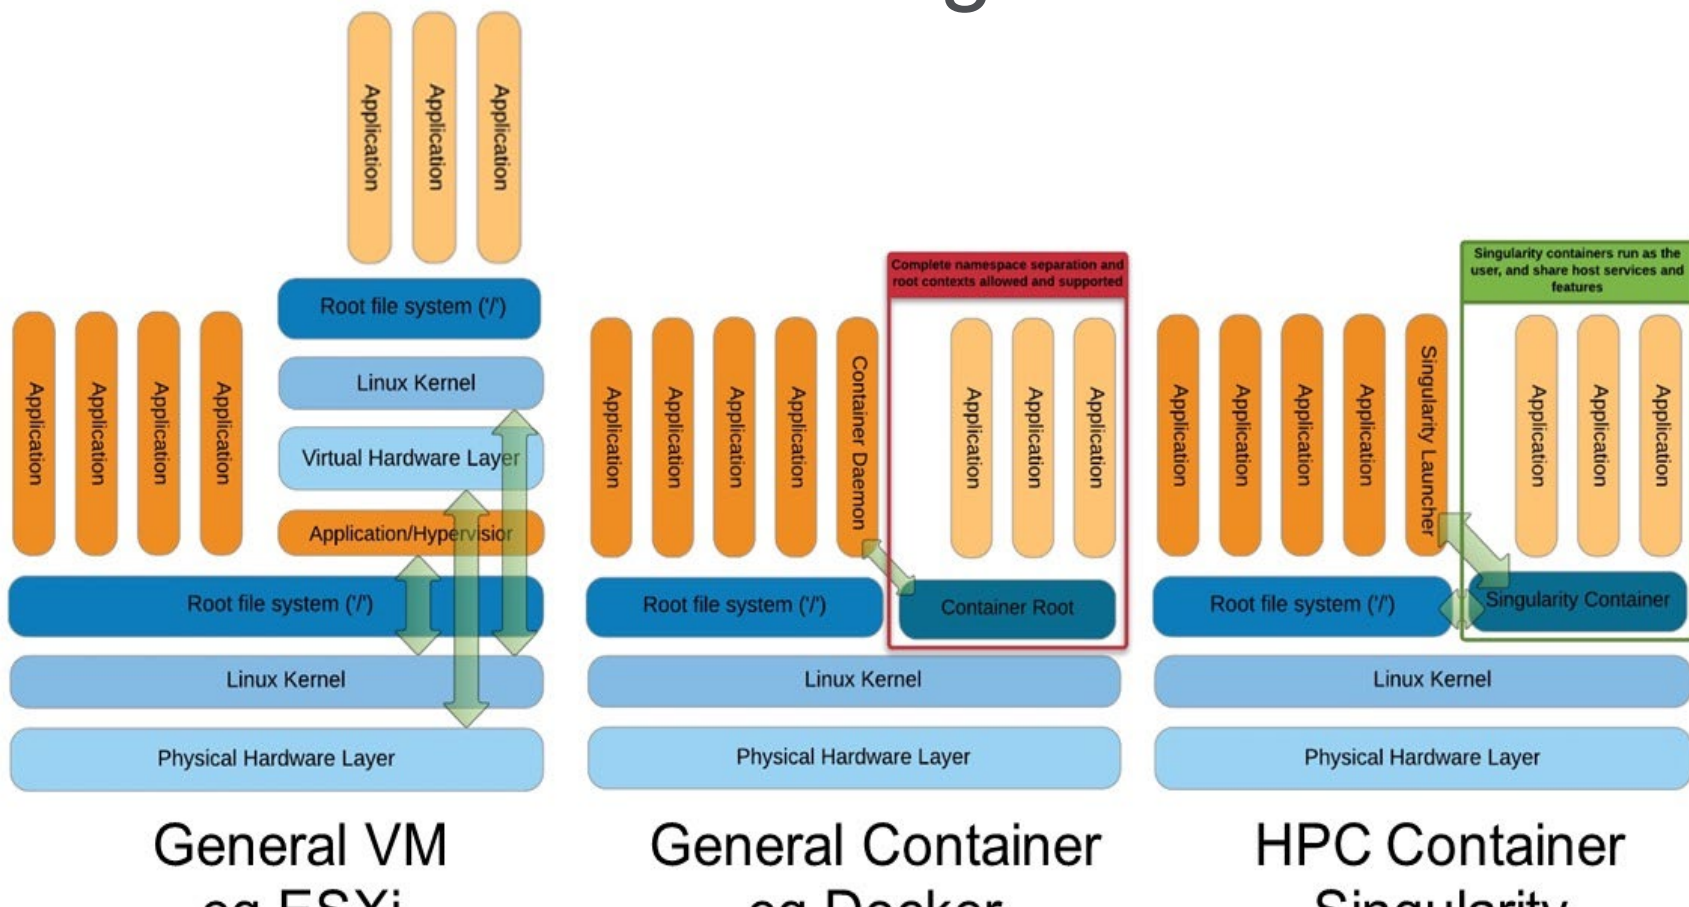
Using a prebuilt image from a registry?

- Decide whether to build your own
- Look at the layers to see what is excess
- Build files on GitHub can reveal unnecessary additions

Building your own is the best way to ensure you have only what you need

Lifecycle of a Container

Data Management



Lifecycle of a Container

Data Management

Manage data in Docker:

<https://docs.docker.com/storage/>

- Volumes
<https://docs.docker.com/storage/volumes>
- Bind mounts
<https://docs.docker.com/storage/bind-mounts/>
- tmpfs mounts
<https://docs.docker.com/storage/tmpfs/>

Manage data in Singularity

- Access to the root filesystem:
https://sylabs.io/guides/3.5/user-guide/quick_start.html#working-with-files
- Bind Paths and Mounts:
https://sylabs.io/guides/3.5/user-guide/bind_paths_and_mounts.html
- Persistent Overlays:
https://sylabs.io/guides/3.5/user-guide/persistent_overlays.html

Deploying a Container

Uploading Containers

Containers on public registries can be a great way to

- Share scientific software portably
- Help others reproduce your research
- Create a community around a software/workflow
- Set up automated builds (connect your GitHub repo)

Common Public Registries

- DockerHub
- Sylabs Cloud Hosting (11GB limit)

Deploying a Container

Best Practices for Uploading Containers


Don't upload

- Private data – very important for research
- Private or licensed software

Do include

- Software licenses
- Documentation
- Software and dependencies
- Runscripts for production

Deploying a Container In the Cloud



Will it work in the cloud?


- Moving from HPC adds complexity
 - MPI
 - May require container orchestration
- Data management

Use Docker

- Public cloud providers offer managed services
- Container Orchestration options
- Ease of use



Deploying a Container In the Cloud



What is the right cloud?

- Private clouds using OpenStack have various options
- Public clouds have vendor-specific options galore
- Multi-cloud or generalized tools also available

Security is a consideration



Deploying a Container On HPC Resources



Containers can simplify getting started

- No need to install to your home directory
- No need to pester sysadmins to install your software

Using Singularity on XSEDE

- It's available and secure
- Automatic mounts for easy data access
- Multinode orchestration via Slurm

Job scripts and bind mounts may vary on different systems



Deploying a Container On HPC Resources



Singularity with MPI

- MPI **major version** in the container *must* match the host
<https://sylabs.io/guides/3.5/user-guide/mpi.html>
- Singlenode will work if MPI is run inside the container
- Multinode is more dependent on system
 - Choice of MPI library may be limited by hardware
 - Scheduler integrations, such as Slurm



Security

Root Access

- Sensitive and shared systems cannot allow root access
- XSEDE systems use Singularity
- Another option is Docker Rootless Mode
 - Docker Docs on Rootless Mode
 - DockerCon 2020 Talk on Rootless Mode
 - Recently became a fully supported feature
- Setup a user or users for shared Docker containers (same as shared system)

https://docs.docker.com/develop/develop-images/dockerfile_best-practices/

Security

Cloud VMs

Implement security at a Virtual Machine (VM) level

- Firewall
- Security Groups
- Limit ssh access

For public images

- Pay attention to what they contain
- Look for the Dockerfile
- GitHub repo

Security

Other Security Features

Docker vulnerability scanning

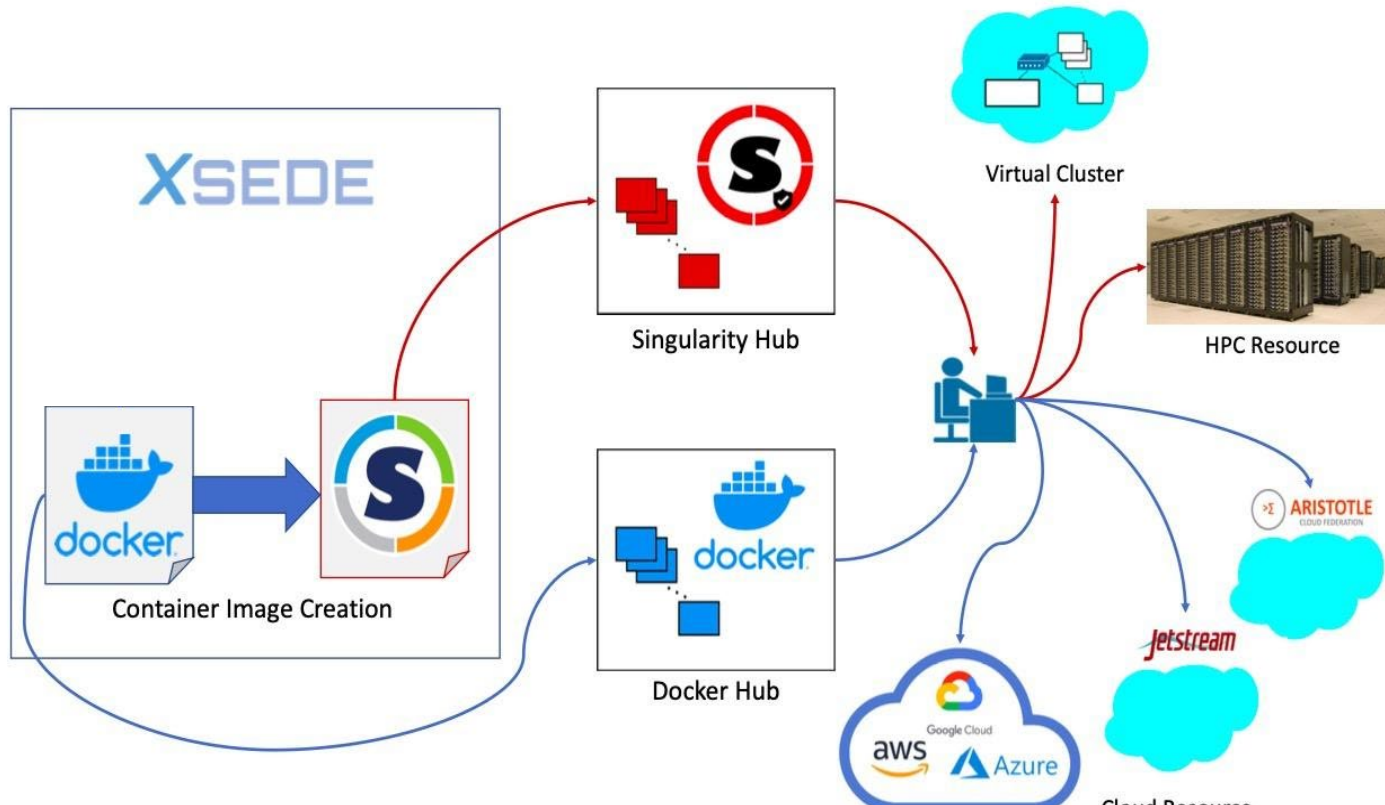
- <https://docs.docker.com/engine/scan>
- <https://github.com/docker/scan-cli-plugin>

Singularity Security Options

- Linux Capabilities
- Encrypted Containers
- https://sylabs.io/guides/3.5/user-guide/security_options.html#security-options


Next Steps

Reproducible Containers



Next Steps


Container Orchestration




+

ANSIBLE

Questions?





Cornell University

XSEDE

Extreme Science and Engineering
Discovery Environment


Thank you!

<https://github.com/XSEDE/Container> Tutorial

Center for Advanced Computing
(CAC) Cornell University

Lifecycle of a Container

Example Container: Radio Astronomy



Lifecycle of a Container

Example Container: Radio Astronomy

Started with a NANOGrav container: `nanograv/nanopulsar`

Based on `jupyter/datascience-notebooks` (includes Python, R, and more)

Wide variety of Radio Astronomy software and tools



Lifecycle of a Container

Example Container: Radio Astronomy

- Started with a NANOGrav container: nanograv/nanopulsar
 - Based on jupyter/datascience-notebooks (includes Python, R, and more)
 - Wide variety of Radio Astronomy software and tools
- After 1 year, needed to be updated: federatedcloud/nanopulsar
- Used for development with additions: federatedcloud/modulation_index
 - ~11GB for just dependencies



Lifecycle of a Container

Example Container: Radio Astronomy

- Started with a NANOGrav container: nanograv/nanopulsar
 - Based on jupyter/datascience-notebooks (includes Python, R, and more)
 - Wide variety of Radio Astronomy software and tools
- After 1 year, needed to be updated: federatedcloud/nanopulsar
- Used for development with additions: federatedcloud/modulation_index
 - ~11GB for just dependencies
- Created a minimal container for production runs
 - ~3GB for just dependencies
 - Docker version: federatedcloud/docker-PRESTO
 - Singularity version: federatedcloud/singularity-PRESTO

